

# C

TA: Sasha Singh

# Some Basics


- C is:
  - **Weakly typed** -you can cast anything into anything
  - **Pass by value** - if you pass something into a function, it will be copied
    - For passing references, use pointers (more into that later)
  - **Compiled**
    - Meaning significantly better runtime
  - Sometimes annoying :(

# Variable and Pointer in C

- Memory is storage with address labels
- Explain the following expressions in C:
  - **int x;**
    - Variable **x** holds 4 bytes, representing an integer
  - **int\* z;**
    - Variable **z** holds 4 bytes, representing a memory location
      - **Memory location = address = pointer**
    - This location holds an integer
  - **int a = \*z;**
    - z holds a location
    - **\*z** gets what's in the location (an integer)
    - Put the integer value in variable a

# Pointer Diagram


- **int x = 61;**
- **int \*y = &x**
- **int \*\*z = &y**
- What are the values of the following?:
  - X is equal to... **61**
  - &x is equal to... **0x0004**
  - Y is equal to... **0x0004**
  - &y is equal to... **0x0008**
  - \*y is equal to... **61**
  - \*\*z is equal to... **61**



# Arrays and Pointer Math

Each int = 4 bytes



- Arrays are pointers
  - Where do we point? The first element in the array
- `int arr[4] = {2, 4, 6, 8};`
- Index using brackets, `a[1] == 4`
- What is the value of the following expressions?
  - **Arr**
    - `0x0004`
  - **\*arr**
    - `2`
  - **Arr + 1**
    - `0x0008`
    - `Add arr + sizeof(int) ⇒ arr + 4 ⇒ 0x0008`
  - **\*(arr + 1)**
    - `4`



# Endianness

Big Endian: write out value towards higher address, the “big end”

Little Endian: write out value towards lower address, the “small end”



# Strings

- Strings in C are **pointers** to 1 or more characters
- The end of a string is given by the null terminator **'\0'**
  - Be careful when you malloc
- `sizeof (char) == 1`
- String lengths can be found with `strlen` (this does NOT include **'\0'**)
  - `int strlen(char *str);`
- String contents can be moved with `strcpy` (DOES copy **'\0'**)
  - `char * strcpy(char *dest, char *src);`
- String contents can be compared with `strcmp`
  - `int strcmp(char *s1, char *s2);`

# Strings Practice

- `char *str = "Hello";`
  - Where is `str[0]` stored?
  - Stored in **static**

**literal "Hello" stored in static, str is just pointing to it's address in static**
- `char b[6] = "Hello";`
  - Where is `b[0]` stored?
  - Initializing an array stores onto the **stack**

**creating space on stack for this variable and then copying chars into variable**
- `char *c = malloc (sizeof (char) * 6);`
- `strcpy (c, "Hello");`
  - Where is `c[0]` stored?
  - Malloced strings go on the **heap**. Make sure to allocate space for the null terminator!

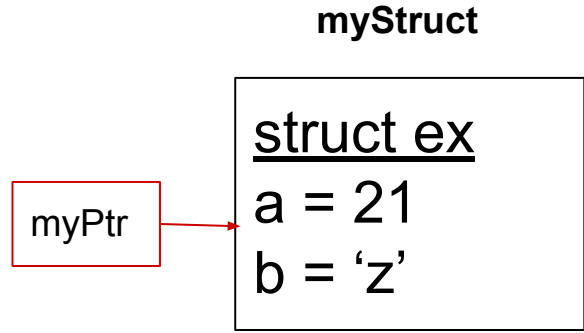
**malloc/calloc/realloc always => heap**



# Structs

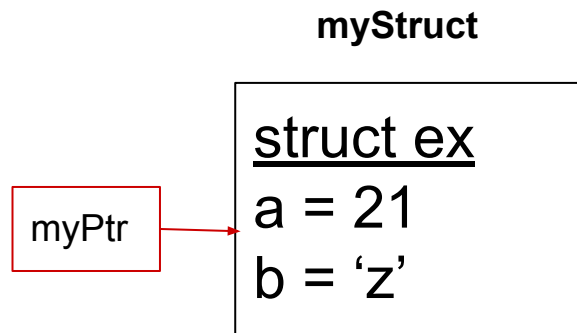
```
struct ex {  
    int a;  
    char b;  
};
```

- Similar to objects in Java, classes in Python
- Struct ex myStruct;
- myStruct.a = 21;
- myStruct.b = 'z'
- myStruct \*myPtr = &myStruct



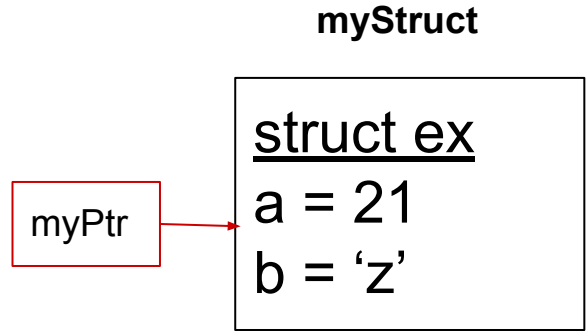
# Careful with Functions!

- Assume we are given myStruct (to the right)
- Say we want a function, **funcA**, that should modify struct.a to be 4
- What is the issue with calling **funcA(myStruct)**?
  - Passing in the struct directly
  - **C is pass by value**



# Careful with Functions!

```
void funcA(struct arg) {  
    arg.a = 4;  
    return;  
}
```


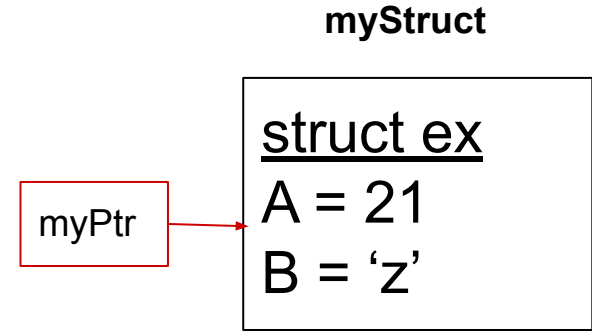


We call funcA(myStruct)!

# Careful with Functions!

```
void funcA(struct arg) {  
    arg.a = 4;  
    return;  
}
```


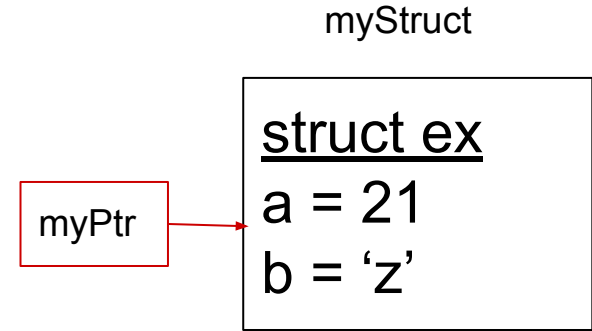
We call funcA(myStruct)!



# Careful with Functions!

```
void funcA(struct arg) {  
    arg.a = 4;  
    return;  
}
```



We call funcA(myStruct)!



# Careful with Functions!

```
void funcA(struct arg) {  
    arg.a = 4;  
    return;  
}
```

The original struct is unchanged!



# Careful with Functions!

- Instead, pass a pointer if you want to retain changes

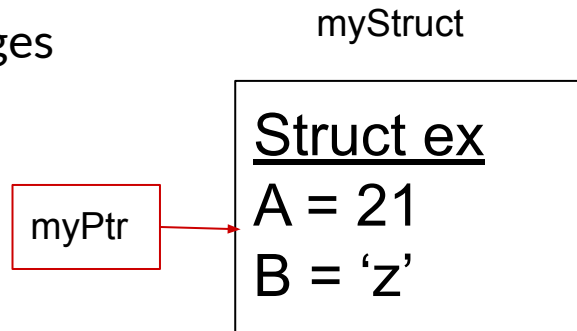
```
Void funcA(Struct *ptr) {
```

```
    (*ptr).A = 4    or    ptr->A = 4
```

```
    Return;
```

```
}
```

- Ptr points to myStruct
- (\*ptr) is mStruct
- (\*ptr).A accesses the field A
  - Dot operator accesses fields
- (\*ptr).A == ptr->A



# Memory Allocation

- `malloc(size_t size)` allocates a block of `size` bytes and returns the start of the block. The time it takes to search for a block is generally not dependent on `size`.
- `calloc(size_t count, size_t size)` allocates a block of `count * size` bytes, sets every value in the block to zero, then returns the start of the block.
- `realloc(void *ptr, size_t size)` "resizes" a previously-allocated block of memory to `size` bytes, returning the start of the resized block.
- `free(void *ptr)` deallocates a block of memory which starts at `ptr` that was previously allocated by the three previous functions.



# C generics

**int\*** pointers only point to **ints**, **char\*** pointers only point to **chars/strings**

**void \* pointers - can point to any type of data**

cannot dereference! **why ?**

**int** \*a; **int** deref\_val = \*a -> how many bytes do we read? **4**

**void**\* b, \*b -> how many bytes do we read? **no clue**

# C generics

if we can't dereference, how do we move data in void pointers around ?

`memcpy` (void\* dest, void\* src, size\_t num\_bytes)


`memmove` (void\* dest, void \*src, size\_t num\_bytes)

move `num_bytes` bytes from the `src` pointer into the `dest` pointer

# Memory Structure

- Memory goes into 1 of 4 segments
  - **Code:**
    - Actual program, RISC-V instructions
    - Does not change in size
  - **Static/Data:**
    - Memory for the whole program
    - Does not change in size
  - **Stack:**
    - Dynamic memory
    - Lives with a function call, local vars
    - Grows Down
  - **Heap:**
    - Dynamic memory
    - Expanded with malloc, shrunk with free
    - Grows Up

$\sim FFFF\ FFF_{hex}$



$\sim 0_{hex}$

# practice ! fa23 midterm

Q2



(C)

(20 points)

Recall in lab, we implemented some functions for a `vector_t` struct. In this question, we will add support for slicing a `vector_t`.

To do so, we've made some updates to the `vector_t` type from lab. You may assume that all necessary standard libraries are included.

```
typedef struct vector_t {
    // Number of elements in the vector; you may assume size > 0
    size_t size;

    // Pointer to the start of the vector
    int* data;

    // Number of child slices
    size_t num_slices;

    // Array of the vector's child slices, or NULL if num_slices == 0
    struct vector_t** slices;

    // true if the vector is a child slice of another vector, otherwise false
    bool is_slice;
} vector_t;
```

Implement the function `vector_slice`, which should return a slice of a `vector_t` at the given indices, with the following signature:

- `vector_t* v`: A pointer to the parent vector to create the slice from.
- `int start_index`: The beginning index of the new slice's data (inclusive)
- `int end_index`: The ending index of the new slice's data (exclusive). You may assume that `end_index > start_index`.
- Return value: A `vector_t*` representing data as described by `start_index` and `end_index`. A parent `vector_t` shares (portions of) its `data` array with all of its descendant slices.

For example:

```
// vec_a has the elements [0, 1, 2, 3, 4]
vector_t* vec_a = /* omitted */;

// vec_b should be of size 2 and have the elements [1, 2]
vector_t* vec_b = vector_slice(vec_a, 1, 3);

vec_b->data[1] = 10;
// At this point, vec_a should be [0, 1, 10, 3, 4]
//                      vec_b should be [1, 10]
```

```

1 vector_t* vector_slice(vector_t* v, int start_index, int end_index) {

2     vector_t* slice = _____;
                                   Q2.1

3     if (slice == NULL) { allocation_failed(); }

4     slice->_____
                                   Q2.2

5     slice->_____
                                   Q2.3

6     slice->_____
                                   Q2.4

7     v->slices = _____;
                                   Q2.5

8     if (v->slices == NULL) { allocation_failed(); }

9     v->slices[_____] = _____;
               Q2.6                      Q2.7

10    v->_____
               Q2.8

11    return slice;

12 }

```

```

1 vector_t* vector_slice(vector_t* v, int start_index, int end_index) {
2     vector_t* slice = calloc(1, sizeof(vector_t));
3     if (slice == NULL) { Q2.1allocation_failed(); }
4     slice->size = end_index - start_index;
5     slice->data = &v->data[start_index]; Q2.2
6     slice->is_slice = true; Q2.3
7     v->slices = realloc(v->slices, (v->num_slices+1)*sizeof(vector_t*)); Q2.4
8     if (v->slices == NULL) { Q2.5allocation_failed(); }
9     v->slices[v->num_slices] = slice; Q2.6 Q2.7
10    v->num_slices += 1; Q2.8
11    return slice;
12 }

```

To accommodate these changes to the `vector_t` type, we need to update the `vector_delete` function to properly free our new data structure. When a `vector_t` is deleted, all descendant slices of the `vector_t` should also be freed. You may assume that `vector_delete` is only called on a `vector_t` that is not a slice.

```
1 void vector_delete(vector_t* v) {  
2     if (_____) {  
3         _____;  
4     }  
5     for (int i = 0; i < v->num_slices; i++) {  
6         _____;  
7     }  
8     if (v->num_slices > 0) {  
9         _____;  
10    }  
11    _____;  
12 }
```



```
1 void vector_delete(vector_t* v) {
2     if (!v->is_slice) {
3         free(v->data);
4     }
5     for (int i = 0; i < v->num_slices; i++) {
6         vector_delete(v->slices[i]);
7     }
8     if (v->num_slices > 0) {
9         free(v->slices);
10    }
11    free(v);
12 }
```

# Practice

## Spring 2019, Q6A

What lines are  
incorrect, and  
what would you  
change?

### Problem 6 C Reading

(16 points)

The function `parse_message` takes two inputs: an array of strings, and the length of the array. It copies the strings from the input array into a new buffer, ending the buffer with a NULL ptr rather than specifying a size. However if any of the strings are the string "STOP", then it terminates early and returns only strings before the stop message, again ending with a NULL terminator.

- (a) The function below contains *at most 5 bugs* which cause the function to non-deterministically exhibit incorrect behavior. Bubble in the lines of code that may produce errors. **You may select more than one line.**

You may assume all calls to `malloc` succeed, `arr` and its contents are never NULL, `arr` always has at least `size` allocated, and we are using C99.

```
☐ 1. char** parse_message (char** arr, size_t size) {  
☐ 2.     int init_size = 8;  
☐ 3.     char **output = malloc (sizeof (char *) * init_size);  
☐ 4.     int i;  
☐ 5.     for (i = 0; i < size; i++) {  
☐ 6.         char *pointer = * arr + i;  
☐ 7.         if (pointer == "STOP") {  
☐ 8.             break;  
☐ 9.         } else if (init_size == i - 1) {  
☐ 10.             init_size *= 2;  
☐ 11.             realloc (output, sizeof (char *) * init_size);  
☐ 12.         }  
☐ 13.         output[i] = malloc (sizeof (char) * strlen (pointer));  
☐ 14.         strcpy (output[i], pointer);  
☐ 15.     }  
☐ 16.     output[i] = NULL;  
☐ 17.     return output;  
☐ 18. }
```

# Practice

## Spring 2019, Q6A

- 6: missing parentheses around the `arr + i`
- 7: incorrect method of comparison
- 9: we resize too late
- 11: `realloc` returns a new pointer and we don't do anything with it :(
- 13: didn't `malloc` for null terminator

### Problem 6 C Reading

(16 points)

The function `parse_message` takes two inputs: an array of strings, and the length of the array. It copies the strings from the input array into a new buffer, ending the buffer with a NULL ptr rather than specifying a size. However if any of the strings are the string "STOP", then it terminates early and returns only strings before the stop message, again ending with a NULL terminator.

- (a) The function below contains *at most 5 bugs* which cause the function to non-deterministically exhibit incorrect behavior. Bubble in the lines of code that may produce errors. **You may select more than one line.**

You may assume all calls to `malloc` succeed, `arr` and its contents are never NULL, `arr` always has at least `size` allocated, and we are using C99.

```
☐ 1. char** parse_message (char** arr, size_t size) {  
☐ 2.     int init_size = 8;  
☐ 3.     char **output = malloc (sizeof (char *) * init_size);  
☐ 4.     int i;  
☐ 5.     for (i = 0; i < size; i++) {  
☐ 6.         char *pointer = * arr + i;  
☐ 7.         if (pointer == "STOP") {  
☐ 8.             break;  
☐ 9.         } else if (init_size == i - 1) {  
☐ 10.             init_size *= 2;  
☐ 11.             realloc (output, sizeof (char *) * init_size);  
☐ 12.         }  
☐ 13.         output[i] = malloc (sizeof (char) * strlen (pointer));  
☐ 14.         strcpy (output[i], pointer);  
☐ 15.     }  
☐ 16.     output[i] = NULL;  
☐ 17.     return output;  
☐ 18. }
```

```
1 #define MAX_BORROWS 25
2
3 typedef struct {
4     char* book_name;
5     bool borrowed;
6 } Book;
7
8 typedef struct {
9     char* user_id;
10    Book* borrowed_books[MAX_BORROWS];
11 } User;
12
13 typedef struct {
14     User* users;
15     int users_len;
16     Book* books;
17     int books_len;
18 } Library;
```

The city of Eddy B.C wants to build a new library! The `init_users` function receives the following input:

- `Library* lib`: A pointer to an uninitialized `Library` struct. You may assume that memory has already been correctly allocated on the heap for the `Library` struct.
- `char** user_ids`: An array of well-formatted strings of nonzero length except the last element. The last element is `NULL`. You may assume that all strings are allocated on the stack.

The function should make sure the following properties are held:

- `users_len` should be set to the number of strings in `user_ids`.
- Each `User` in `users` should be initialized as follows:
  - The `user_id` of the `i`th `User` in `users` should be set to the `i`th string in `user_ids`.
  - `borrowed_books` should be an array of `NULL`s to indicate that no `Book` has been borrowed.
- Every `User` and its contents must persist through function calls.

(Question 2 continued...)

Useful C function prototypes:

```
void* malloc(size_t size);  
void free(void *ptr);  
void* calloc(size_t num_elements, size_t size);  
void* realloc(void *ptr, size_t size);  
  
size_t strlen(char* s);  
char* strcpy(char* dest, char* src);  
  
// memset sets the first num bytes of the block of memory pointed to by ptr  
// to the specified value (interpreted as an unsigned char).  
void* memset(void* ptr, int value, size_t num);
```

(15 points) Fill in `init_users` so that it matches the described behavior. Assume that all necessary C libraries are included.

```
1 void init_users(Library* lib, char** user_ids) {
2   int i = 0;
3   while ( _____ ) {
4     lib _____ = _____;
5     User* cur_user = _____;
6     cur_user _____ = _____;
7     strcpy(cur_user _____, _____);
8     memset(cur_user _____, _____,
              MAX_BORROWS * _____);
9     i++;
10  }
11  lib _____ = i - 1;
12 }
```

**Solution:**

```
1 void init_users(Library* lib, char** user_ids) {
2     int i = 0;
3     while (user_ids[i] != NULL) {
4         lib->users = realloc(lib->users, sizeof(User) * (i + 1));
5         User* cur_user = &lib->users[i];
6         cur_user->user_id = malloc((strlen(user_ids[i]) + 1) * sizeof(char));
7         strcpy(cur_user->user_id, user_ids[i]);
8         memset(cur_user->borrowed_books, 0,
                MAX_BORROWS * sizeof(Book *));
```

(Question 2 continued...)

```
9         i++;
10    }
11    lib->users_len = i - 1;
12 }
```