



UC Berkeley  
Teaching Professor  
Dan Garcia

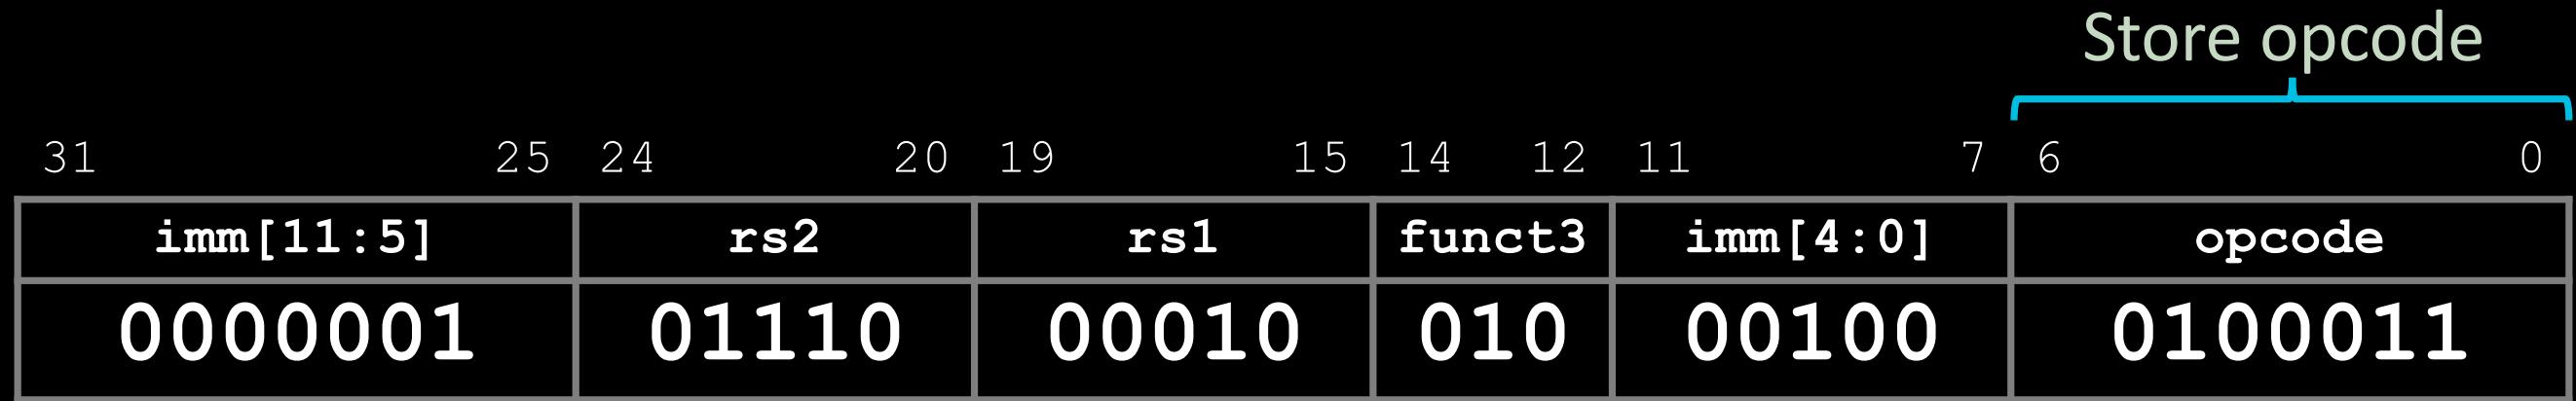
# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)

### RISC-V Instruction Formats, Part II

# Disassembling the S-Format

- What is the corresponding assembly instruction?

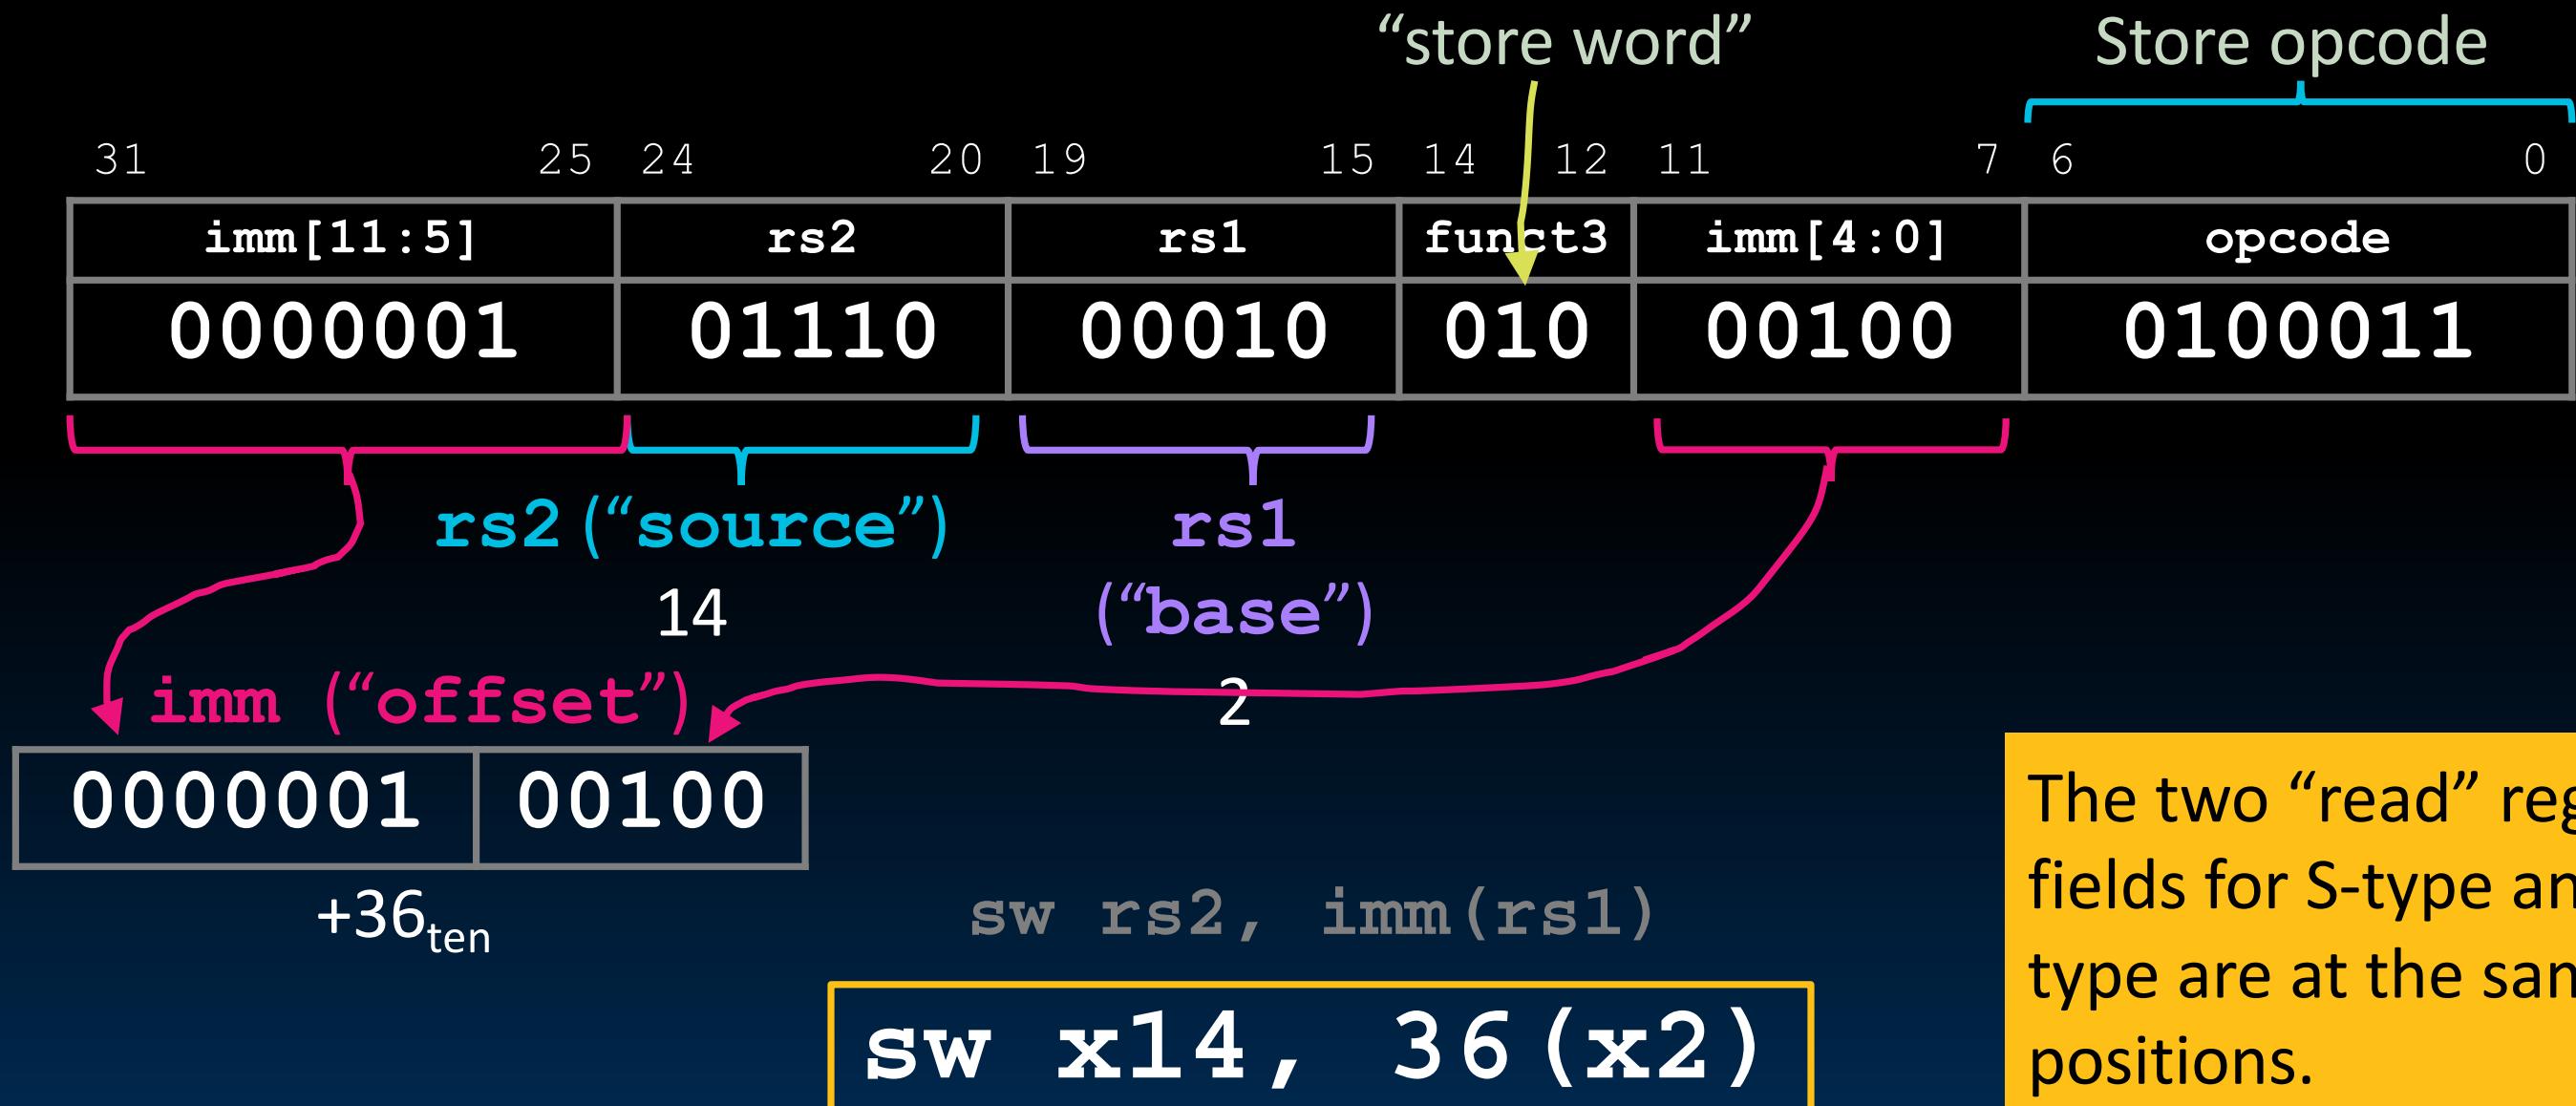


S-Format Lookup table:

			funct3	opcode		
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	sb
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	sh
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	sw

# Disassembling the S-Format

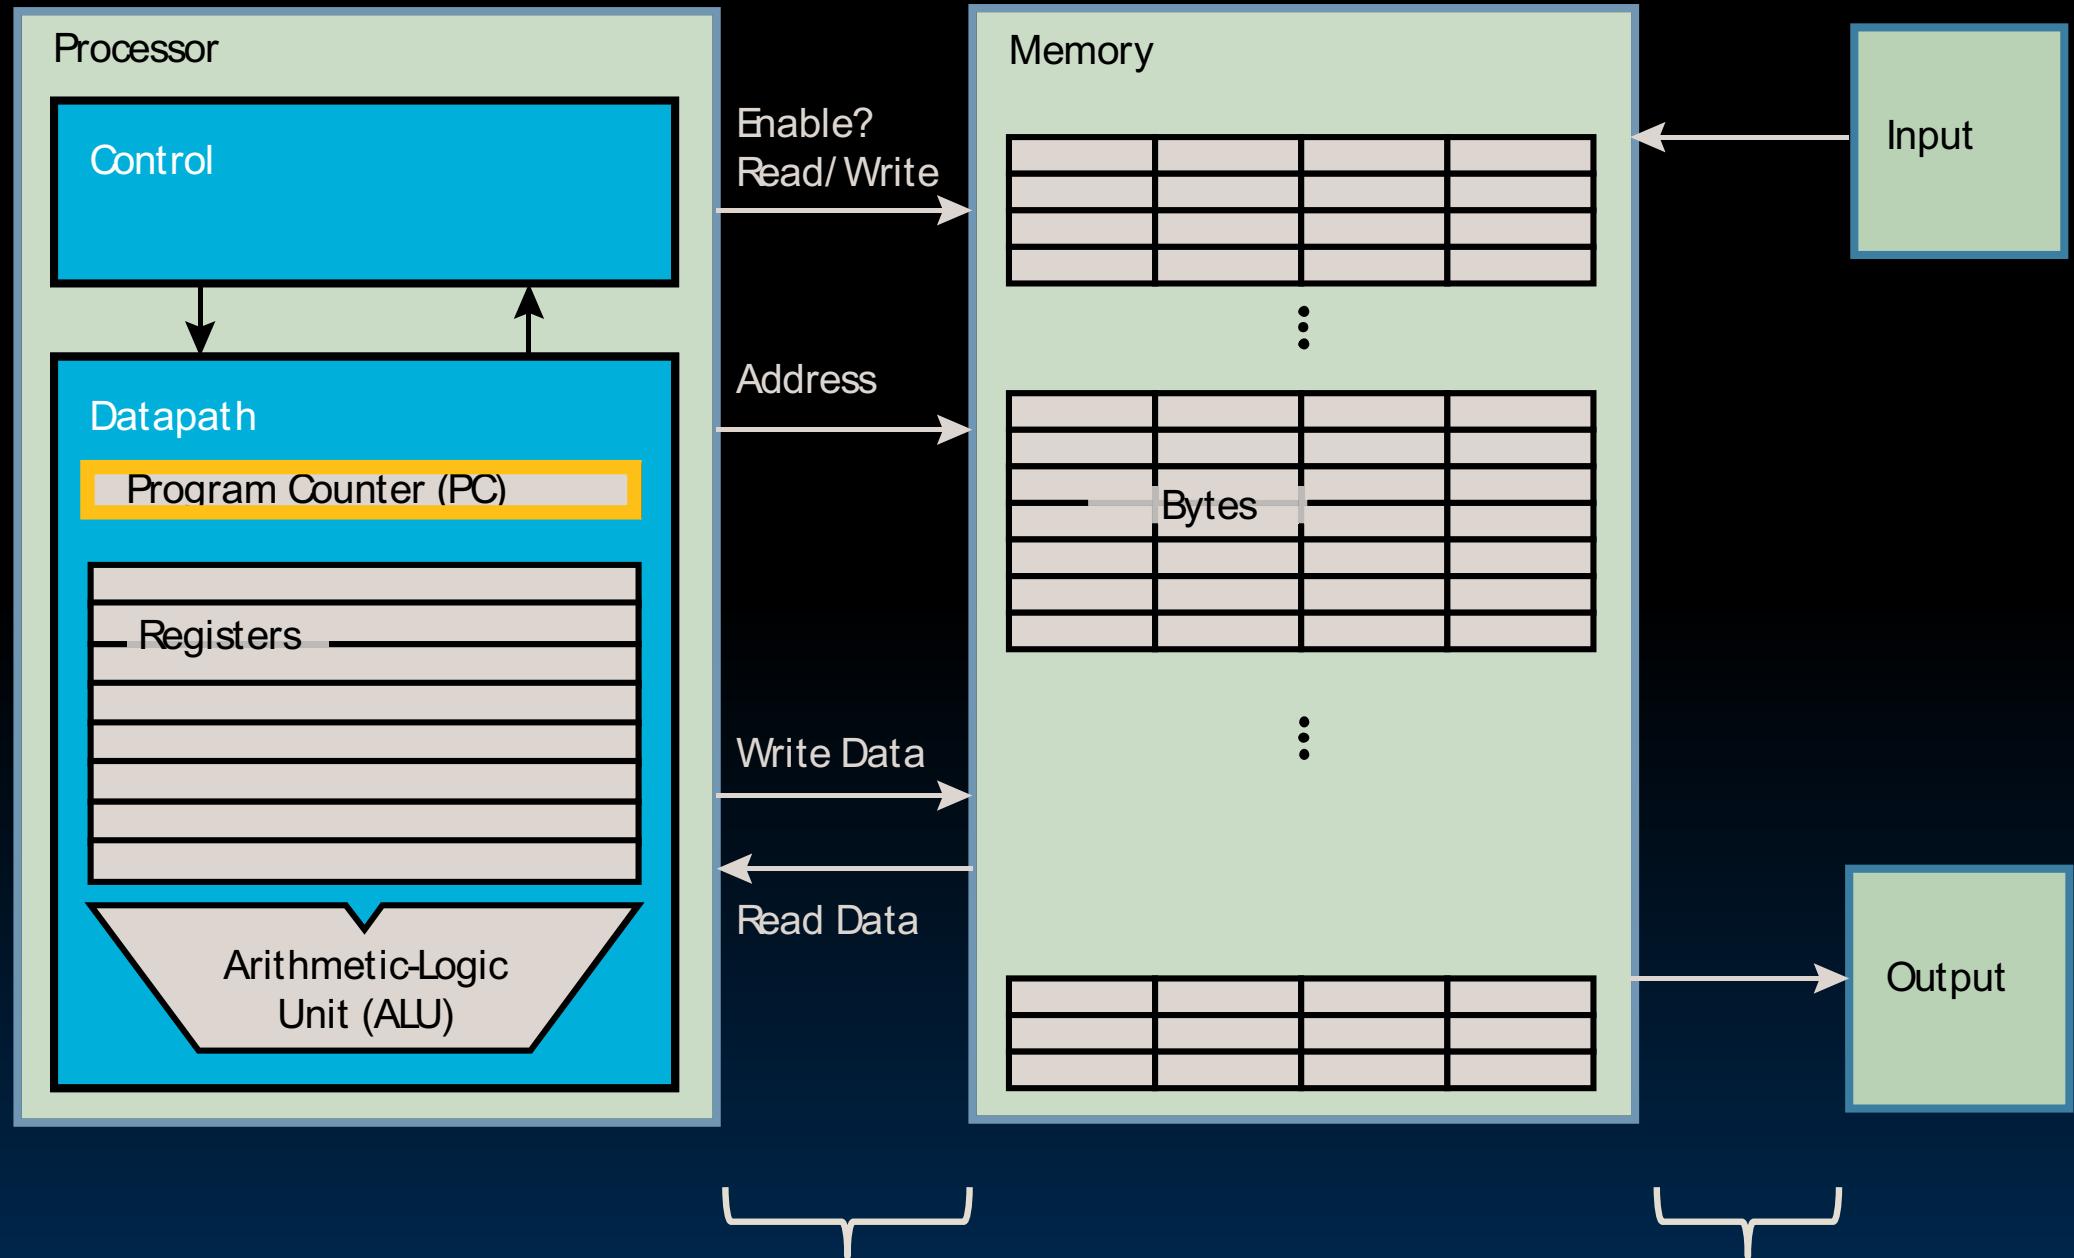
- What is the corresponding assembly instruction?



# Updating the Program Counter (PC)

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- **jalr**: I-Format

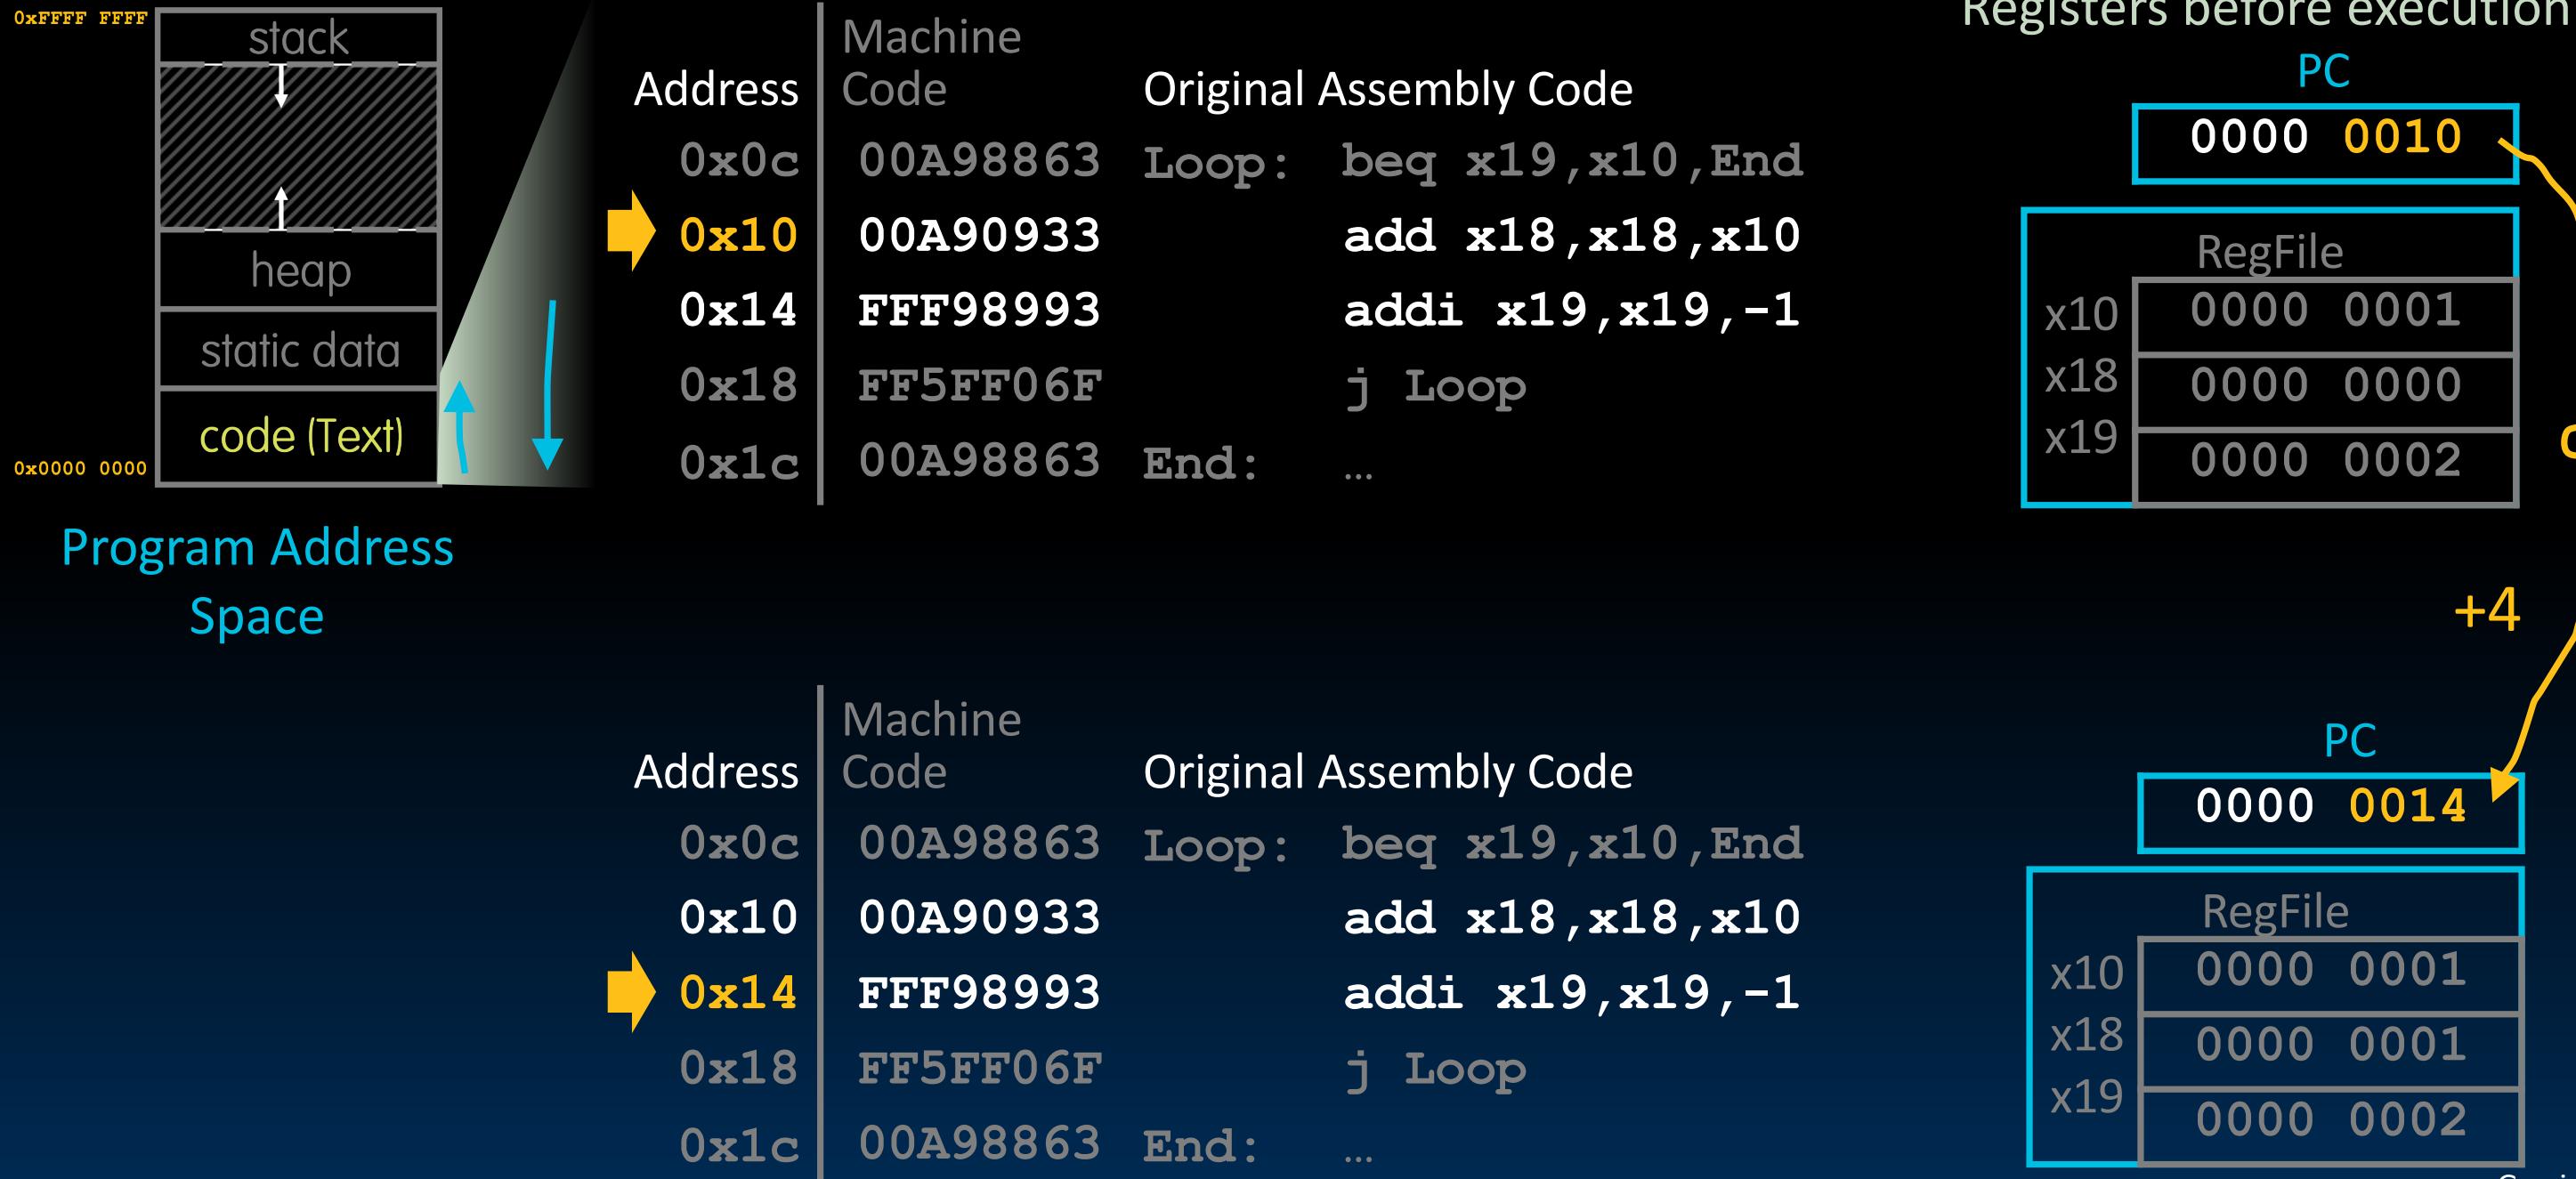
# The Program Counter (PC)



**The Program Counter (PC) is a special internal register inside processor holding byte address of the instruction being executed.**

**It is separate from the 32-registers Register File!**

# $PC = PC + 4$ for most instructions



# Branches Update PC to "Jump"

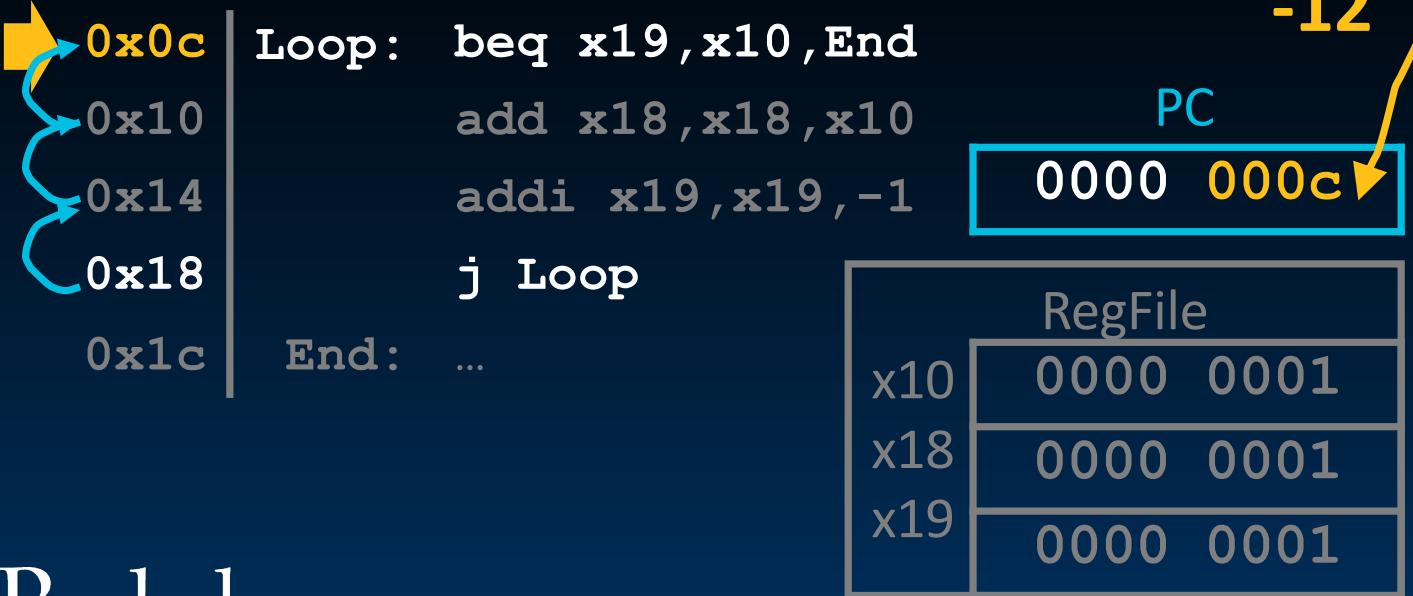
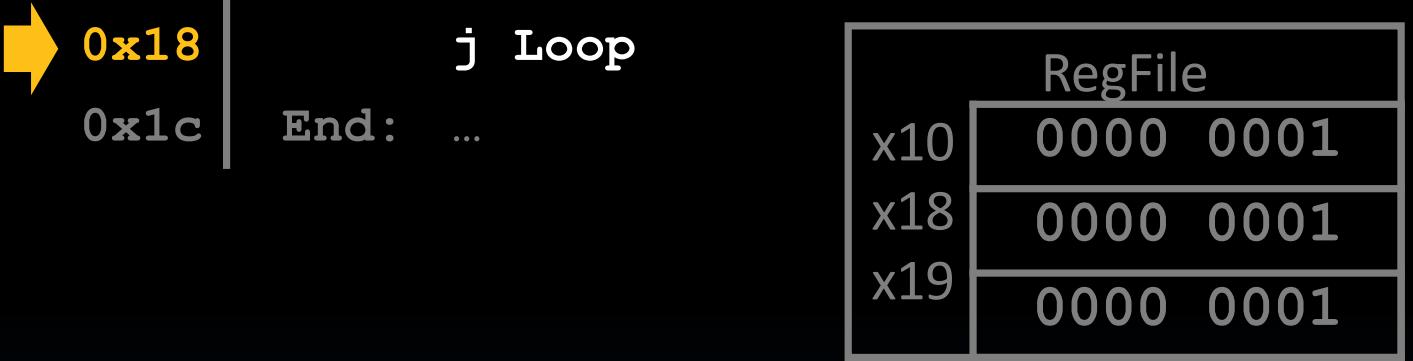
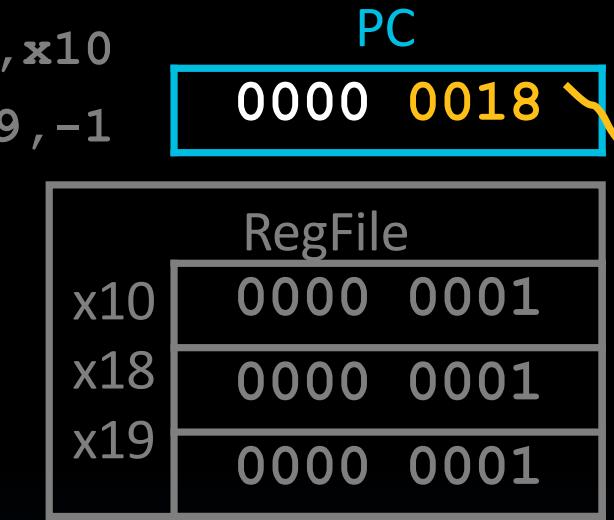


Kris Kross - Jump (Official Video) -

## Unconditional Branches

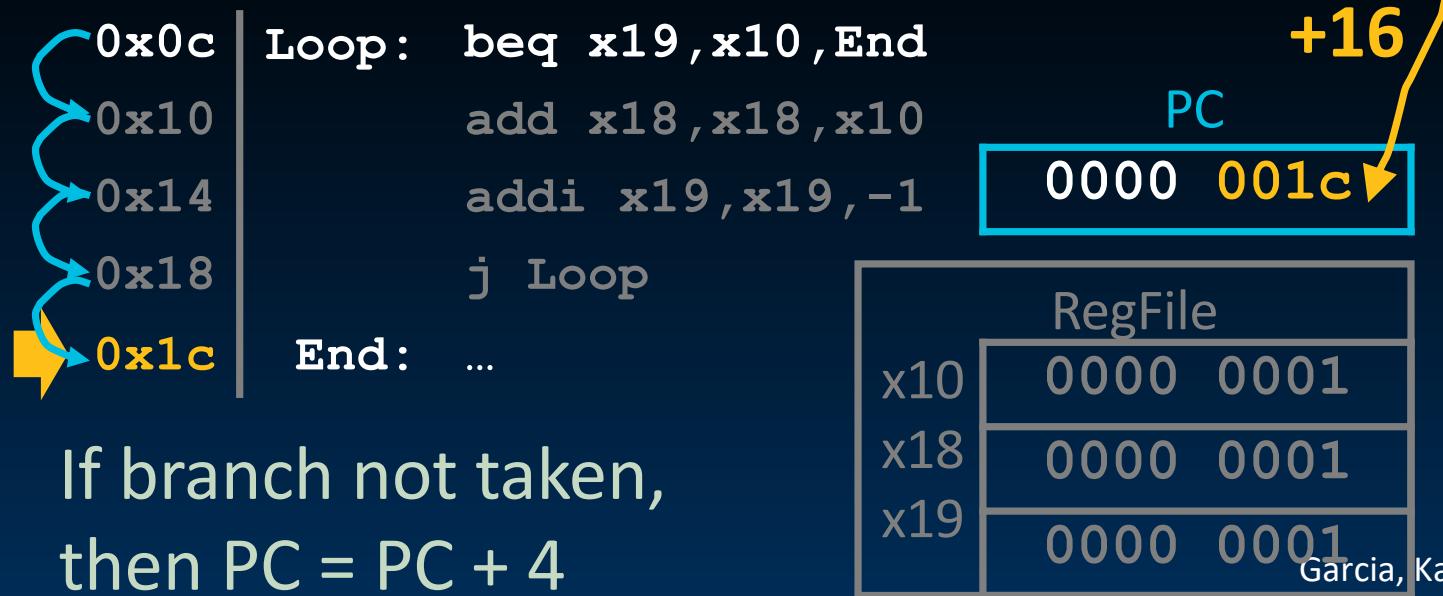
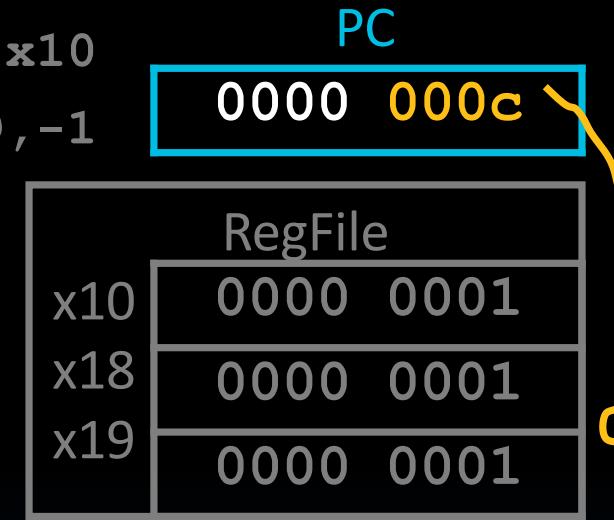
```

0x0c | Loop: beq x19,x10,End
      add x18 x18,x10
      addi x19,x19,-1
      j Loop
      End: ...
  
```



## Conditional Branch (if branch taken)

→ 0x0c | Loop: beq x19,x10,End  
 add x18 x18,x10  
 addi x19,x19,-1  
 j Loop  
 End: ...



If branch not taken,  
then  $PC = PC + 4$

# PC-Relative Addressing

- PC-Relative Addressing: Supply a signed offset to update PC.

$$\text{PC} = \text{PC} + \text{byte\_offset}$$

*“Position-Independent Code”*: If all of code moves, *relative offsets* don’t change!

0x0c	Loop: <code>beq x19, x10, End</code>
0x10	<code>add x18, x18, x10</code>
0x14	<code>addi x19, x19, -1</code>
0x18	<code>j Loop</code>
0x1c	End: ...

Take branch:  
 $\text{PC} = \text{PC} + 16$   
Don’t take branch:  
 $\text{PC} = \text{PC} + 4$

0x400	Loop: <code>beq x19, x10, End</code>
0x404	<code>add x18, x18, x10</code>
0x408	<code>addi x19, x19, -1</code>
0x40c	<code>j Loop</code>
0x410	End: ...

Branches generally change the PC by a small amount, therefore in RISC-V instructions we encode relative offsets as *signed immediates*.

- Contrast with: Absolute Addressing

Supply new address to overwrite PC.  $\text{PC} = \text{new\_address}$

Use sparingly: Brittle to code movement/need to build 32-bit immediate  
(more later)

# B-Format Layout

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- **jalr**: I-Format

# RISC-V Conditional Branches

- Example: **beq rs1, rs2, Label**
- Conditional branches need:
  - Two **read** registers, **no destination register** (like S-Format)
  - Some way to encode label (i.e., where to jump to)
- RISC-V uses **PC-relative addressing to encode labels**.
  - If we *don't* branch:  $\text{PC} = \text{PC} + 4 \text{ bytes}$
  - If we *do* branch:  $\text{PC} = \text{PC} + (\text{relative byte offset to Label})$
- B-Format Fields:
  - 7b (**opcode**) + 3b (**funct3**) + 5b (**rs1**) + 5b (**rs2**) = 20b
  - 12b **immediate field** to represent PC's relative offset.
    - Signed two's complement means we can represent  $\pm 2^{11}$  "units" from the PC.

# Scaling the Offset to Maximize Branch Range

The 12-bit immediate field for conditional branches is:

A. In units of 1 byte.

One branch reaches  $\pm 2^{11}$  bytes from PC  $\rightarrow \pm 2^9 \times 32$

Never branch into middle of instruction.

B. In units of 2 bytes (16-bit “half-words”).

One branch reaches  $\pm 2^{10} \times 32$

Multiply the offset by 2 before adding to the PC.

C. In units of 4 bytes (32-bit words).

One branch reaches  $\pm 2^{11} \times 32$

Multiply the offset by 4 before adding to the PC.



## Scaling the Offset to Maximize Branch Range

The 12-bit immediate field for conditional branches is:

**A. In units of 1 byte.**

- One branch reaches  $\pm 2^{11}$  bytes from PC  $\rightarrow \pm 2^9 \times 32\text{-bit instructions}$ .
- Never branch into middle of instruction.

**B. In units of 2 bytes.**

- One branch reaches  $\pm 2^{10} \times 32\text{-b instructions}$  from PC.
- Multiply the offset by 2 bytes before adding to the PC.

**C. In units of 4 bytes.**

- One branch reaches  $\pm 2^{11} \times 32\text{-b instructions}$  from PC.
- Multiply the offset by 4 bytes before adding to the PC.

(A) A

(B) B

(C) C

(D) None of the above

0%

0%

0%

0%

# Scaling the Offset to Maximize Branch Range

The 12-bit immediate field for conditional branches is:

- X In units of 1 byte.

One branch reaches  $\pm 2^{11}$  bytes from PC  $\rightarrow \pm 2^9 \times 32$

Never branch into middle of instruction.

*Wasteful!* Lower 2 bits would always be wasted (always 00).

- ✓ In units of 2 bytes (16-bit “half-words”).

One branch reaches  $\pm 2^{10} \times 32$

Multiply the offset by 2 before adding to the PC.

- X In units of 4 bytes (32-bit words).

One branch reaches  $\pm 2^{11} \times 32$

Multiply the offset by 4 before adding to the PC.

Does not support 16-bit compressed instructions!

Branch immediate range [-2048, +2047] represents PC relative offsets [-4096, +4094] in 2-byte increments.

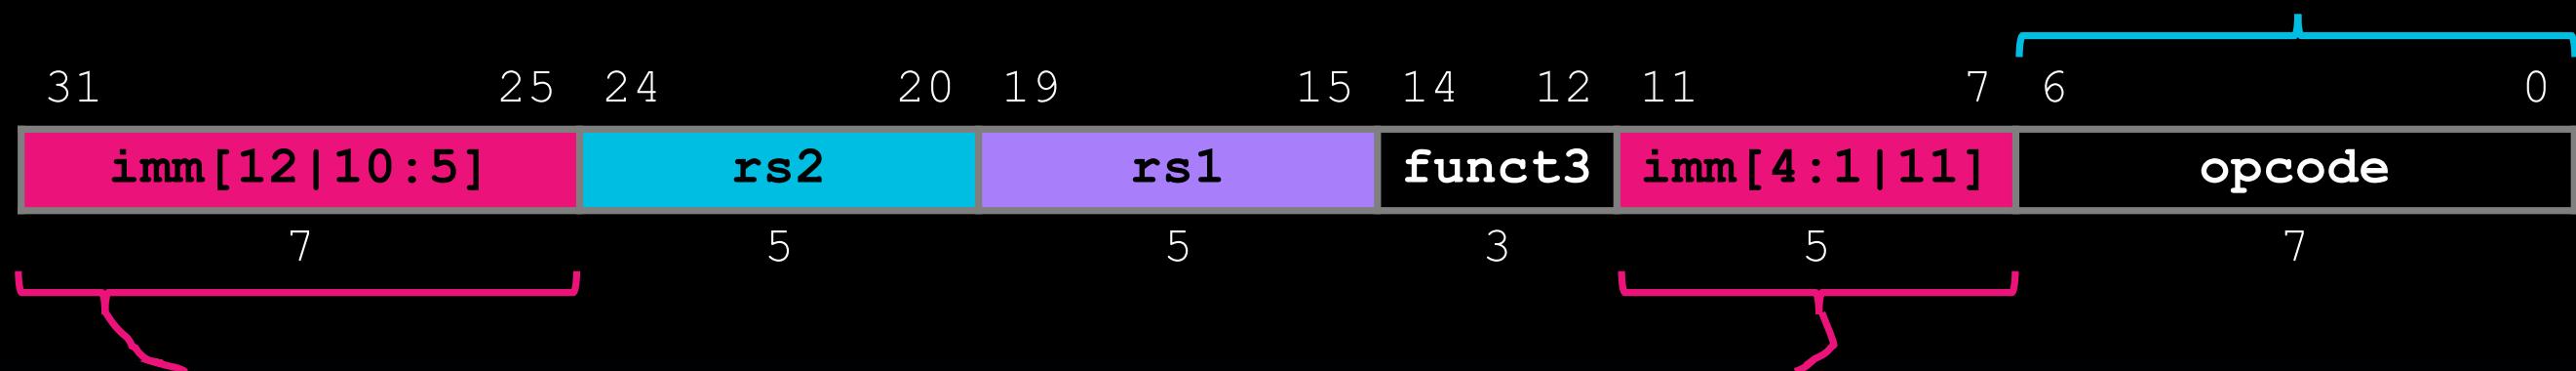
# RISC-V Feature: 16-bit instructions

- RISC-V Base ISA for RV32, RV64, RV128 all have 32-bit wide instructions, but it supports extensions:
  - 16b compressed instructions
  - Variable-length instructions that are multiples of 16b in length
- RISC-V therefore *scales the branch offset by 2 bytes*, even when there are no 16b instructions!
- Implications for this class (which uses RISC-V processors that only support 32-bit instructions):
  - Half of the possible branch targets will be errors.
  - RISC-V Conditional Branches can only reach  $\pm 2^{10}$  32-bit instructions on either side of PC.

# B-Format Instruction Layout

- B-Format (textbook: SB-Type) close to S-Format:  
**opname      rs1, rs2 , Label**

All *conditional* branch instructions have opcode 1100011.



Immediate represents relative offset in increments of 2 bytes (“half-words”)

To compute new PC:  $PC = PC + \text{byte\_offset}$ .

12 immediate bits imply  $\pm 2^{10}$  32-bit instructions reachable:

1 bit: 2's complement (allow  $+\/-$  offset)

1 bit: half-word/16-b instruction support

(Lowest bit of offset is always zero, so no need to store in instruction.)

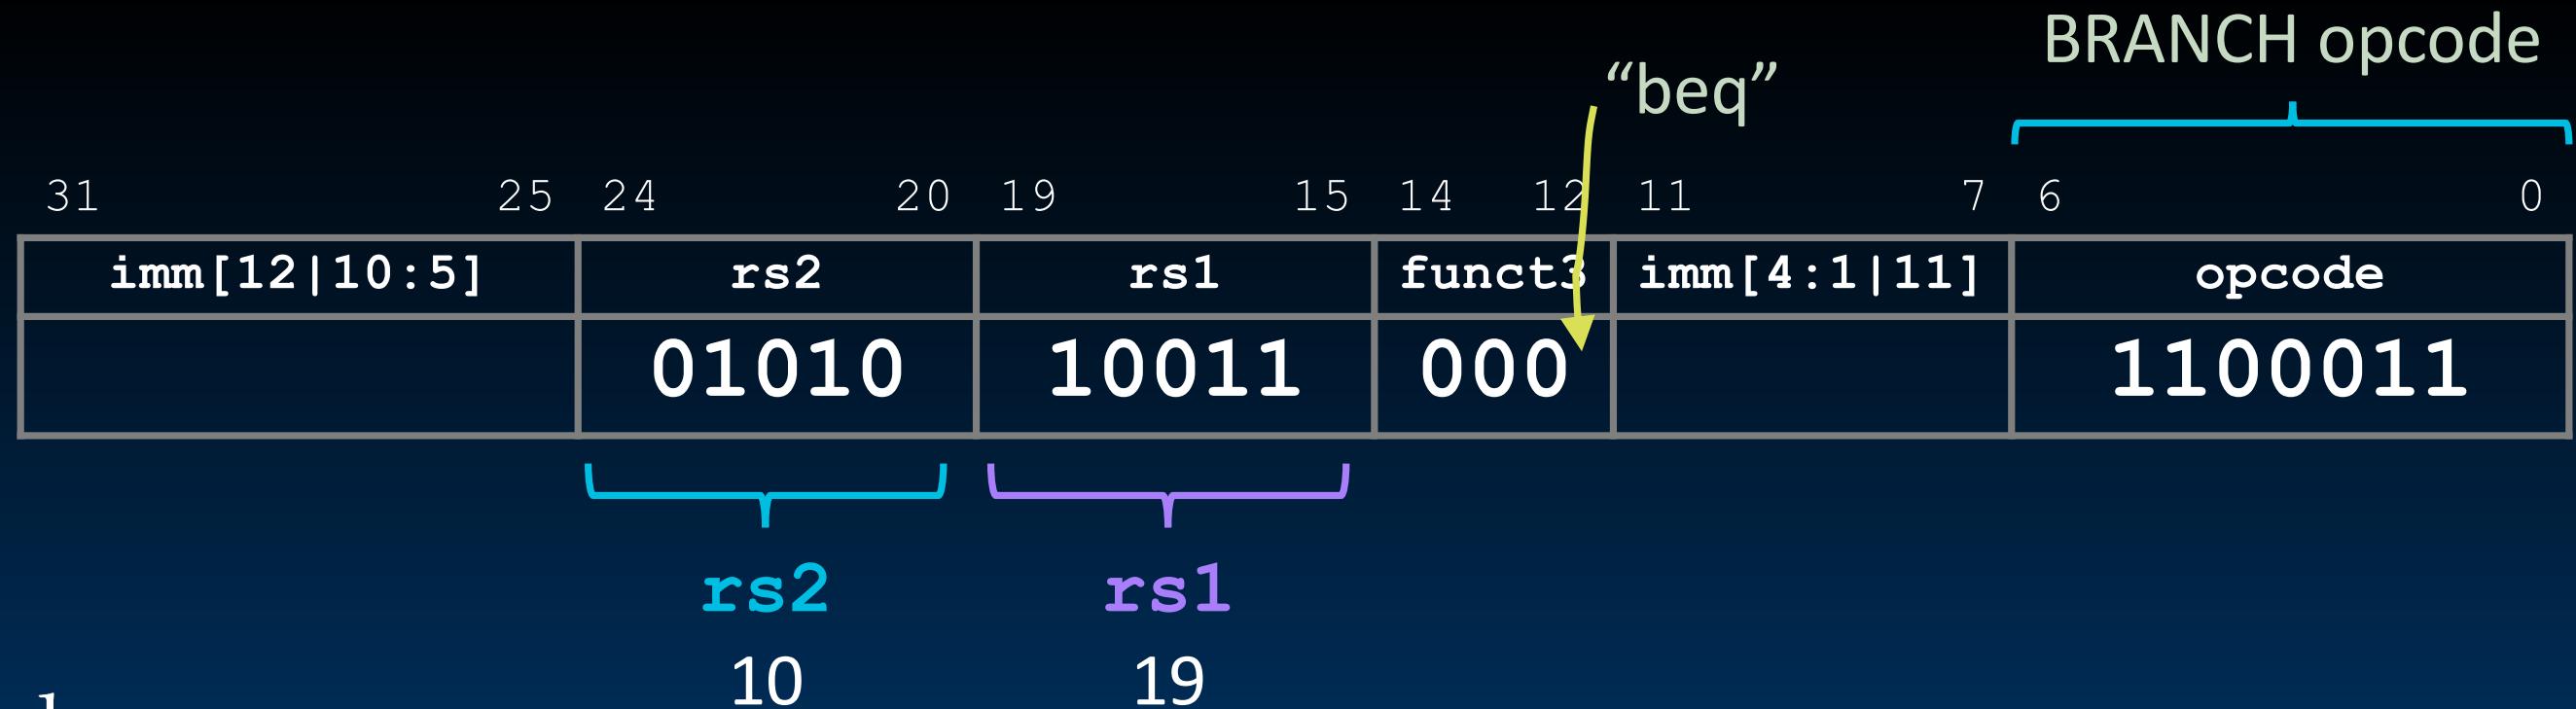


# B-Format Example (1/2)

```
Loop:  beq x19,x10,End  
       add x18,x18,x10  
       addi x19,x19,-1  
       j Loop  
  
End:   ... # target instr
```

How do we encode the  
beq instruction in this  
RISC-V code?

1. Fill in fields (other than immediate).



# B-Format Example (2/2)

+4  
instruct-  
ions {  
Loop:   **beq x19,x10,End**  
          add x18,x18,x10  
          addi x19,x19,-1  
          j Loop  
End:    ... # target instr

How do we encode the  
**beq** instruction in this  
RISC-V code?

2. Construct branch offset (in bytes), then encode immediate.

31	25 24	20 19	15	14	12	11	7	6	0
imm[12 10:5]	rs2	rs1	funct3	imm[4:1 11]	opcode				
0000000	01010	10011	000	10000	1100011				

+4 32-bit instructions = +16 byte offset

12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	0	0	0	0	0	1	0	0	0	0

13-bit signed **byte\_offset** ("imm")

# Why Swirl Immediate Bits Around? HW Design!

- Recall: RISC-V register field positions consistent across instr. formats.
- RISC-V also tries to keep bit positions of immediates consistent:

	31	25	24	20	19	15	14	12	11	7	6	0
I-type	imm[11:5]	imm[4:0]		rs1		funct3		rd			opcode	
	<b>yxxxxxxxxx</b>	<b>wwwwv</b>										
S-type	imm[11:5]		rs2		rs1	funct3	imm[4:0]				opcode	
	<b>yxxxxxxxxx</b>						<b>wwwwv</b>					
B-type	imm[12 10:5]		rs2		rs1	funct3	imm[4:1 11]				opcode	
	<b>zxxxxxxxxx</b>						<b>wwwwy</b>					

Instruction Bit 31 is always the *sign bit* (highest bit to sign extend in immediate)  
 B-type has a *13-bit* immediate encoding b/c of implicit zero in Imm bit 0.

Between S and B, only two bits change meaning:  
 Instruction Bit 7 → S: Imm Bit 0; B: Imm Bit 11  
 Instruction Bit 31 → S: Imm Bit 11; B: Imm Bit 12

RISC-V immediate bit encoding is optimized to reduce hardware cost.

# All six RV32 B-Format Instructions

imm[12 10:5]	rs2	rs1	funct3	imm[4:1 11]	opcode	
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011	b <sub>eq</sub>
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011	b <sub>nne</sub>
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011	b <sub>lt</sub>
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011	b <sub>ge</sub>
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011	b <sub>ltu</sub>
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011	b <sub>geu</sub>

# Conditionally Branching Even Further

- Conditional Branches are typically used for if-else statements, for/while loops.

Usually pretty small (<50 lines of C code)

- B-Format has limited range:

$\pm 2^{10}$  32-bit instructions from current instruction

- What if destination is further away?

Enter the *unconditional jump*!

```
beq x10, x0, far  
# next instr
```

if equal

functionally equivalent

```
bne x10, x0, next  
j far  
next: # next instr
```

if equal

Garcia, Kao



# J-Format Layout

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- **jalr**: I-Format

# The jal instruction and j pseudoinstruction

- **jal** does a **Jump And Link** (i.e., “link and jump”):

**jal rd, Label**

Jump to **Label**

Write address of the *following* instruction to **rd**.

- **Two use cases:**

1. Call a function (and simultaneously save return address in named register)

**jal ra, FuncLabel**

2. Unconditional branch

**j pseudoinstruction: jump and discard return address**

**j Label**



**jal x0, Label**

# The jal instruction and j pseudoinstruction

- jal does a Jump And Link (i.e., “link and jump”):

**jal rd, Label**

Jump to **Label** (i.e., add relative offset to PC)

Write address of the *following* instruction to **rd**.

$$\left. \begin{array}{l} \text{PC} = \text{PC} + \text{offset} \\ \text{rd} = \text{PC} + 4 \end{array} \right\}$$

- Two use cases:

1. Call a function (and simultaneously save return address in named register)

**jal ra, FuncLabel**

2. Unconditional branch

j pseudoinstruction: jump and discard return address

**j Label**

Often used to conditionally branch further than B-Types

**jal x0, Label**



# J-Format Instruction Layout

- J-Format (textbook: UJ-Type) used only for `jal`:  
`jal rd, Label`



Immediate represents relative offset in increments of 2 bytes.

To compute new PC:  $PC = PC + \text{byte\_offset}$ .

20 immediate bits imply  $\pm 2^{18}$  32-bit instructions reachable:

1 bit: 2's complement (allow  $+\/-$  offset), 1 bit: half-word/16-b instruction support

Immediate bit encoding optimized to reduce HW cost

What about jumping further?

Next: How to load 32-bit  
immediate into a register!

# U-Format: Long Immediates

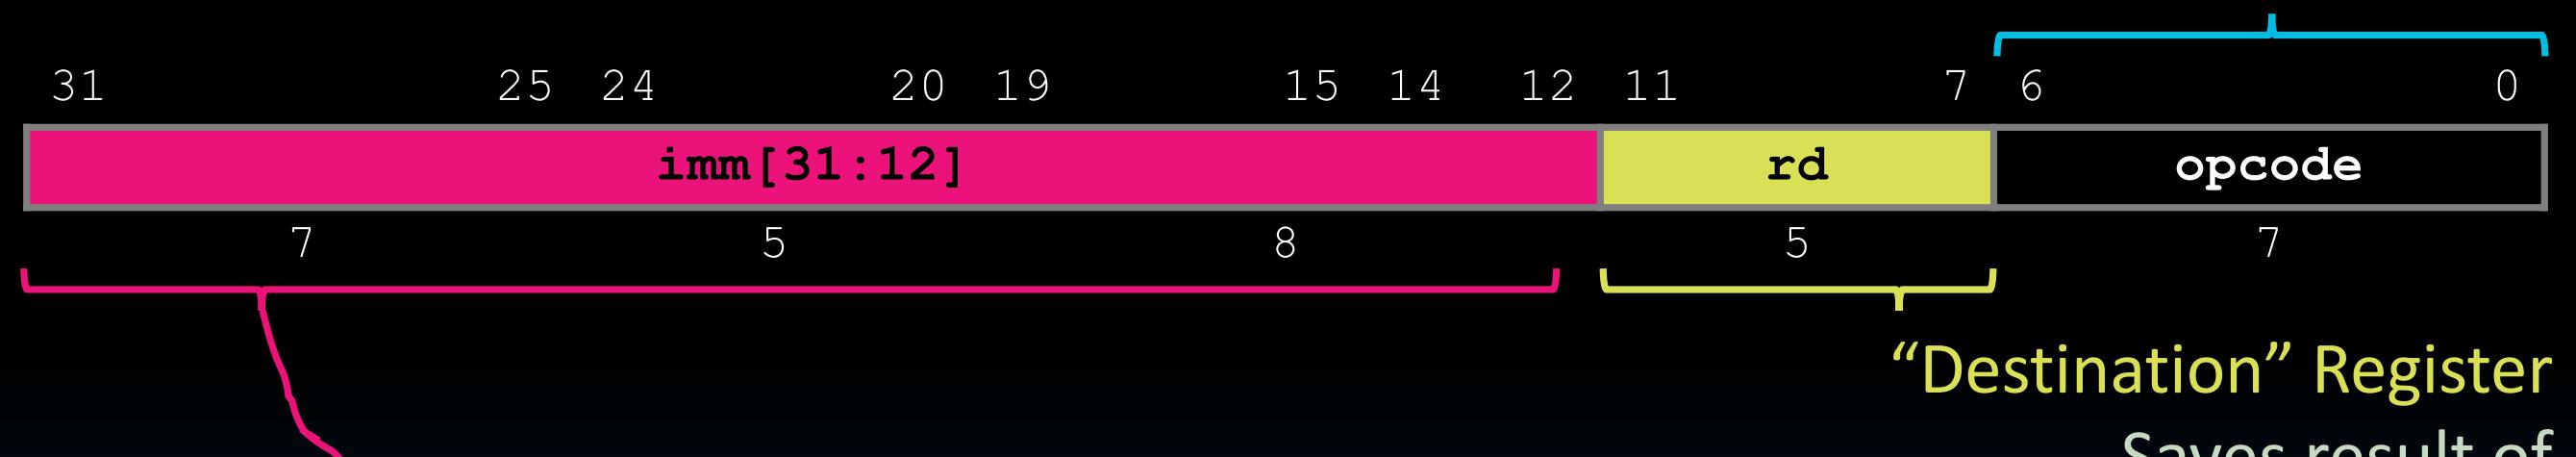
- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- **jalr**: I-Format

# U-Format Instruction Layout

- “Upper Immediate” instructions:

**opname      rd, immed**

U-Format opcodes:  
lui    0110111  
auipc 0010111



Immediate represents *upper 20 bits* of a 32-bit immediate  
operand **imm** = **immed** << 12.



Long Immediate (“**imm**”)

# lui and addi creates Long Immediates

**lui rd, immed**

- The **lui** instruction, Load Upper Immediate:

Write a 20-bit immediate value into  
the upper 20-bits of register rd.

Clear the lower 12 bits.

$$\left. \begin{array}{l} \text{rd} = \text{immed} \ll 12 \\ \end{array} \right\}$$

- lui together with an addi (to set lower 12 bits) can create any 32-bit value in a register:**

```
lui x10, 0x87654      # x10 = 0x87654000
addi x10, x10, 0x321  # x10 = 0x87654321
```

The **li** pseudoinstruction  
(Load Immediate) resolves to  
**lui + addi** as needed,  
e.g., **li x10, 0x87654321**.

# lui Edge Case: addi Extends Sign

- How should we set the immediate 0xB0BACAFE?

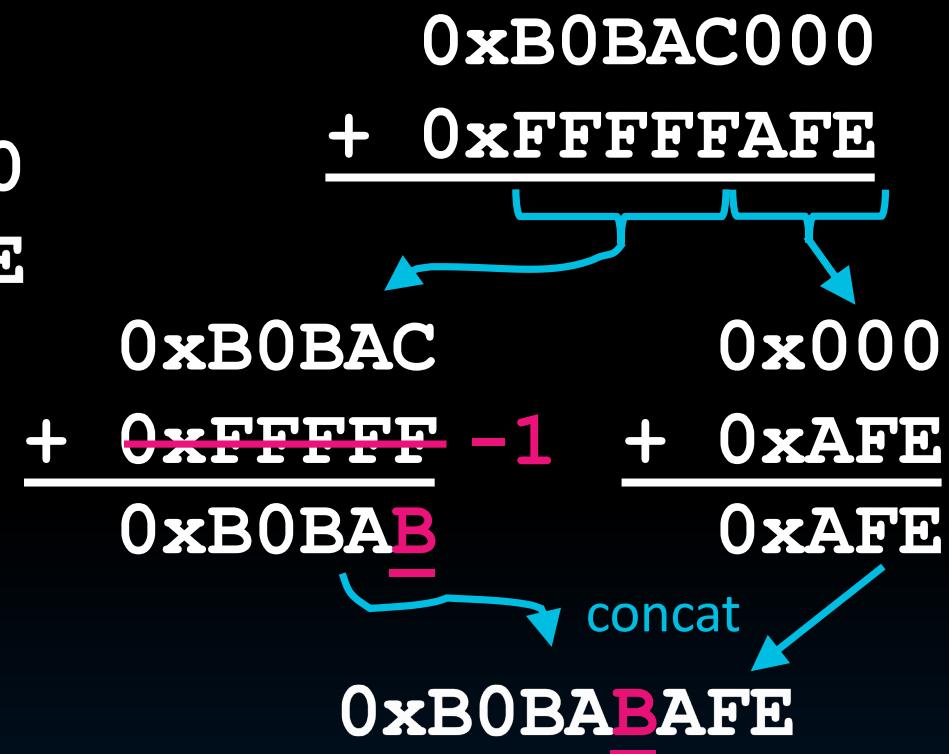


Unfortunately:

```
lui x10, 0xB0BAC      # x10 = 0xB0BAC000  
addi x10, x10, 0xAF   # x10 = 0xB0BABAFE
```

Recall: **addi** sign-extends the 12-bit immediate.

If “sign bit” set, subtracts 1 from the upper 20-bits!



# lui Edge Case: addi Extends Sign

- How should we set the immediate 0xB0BACAFE?

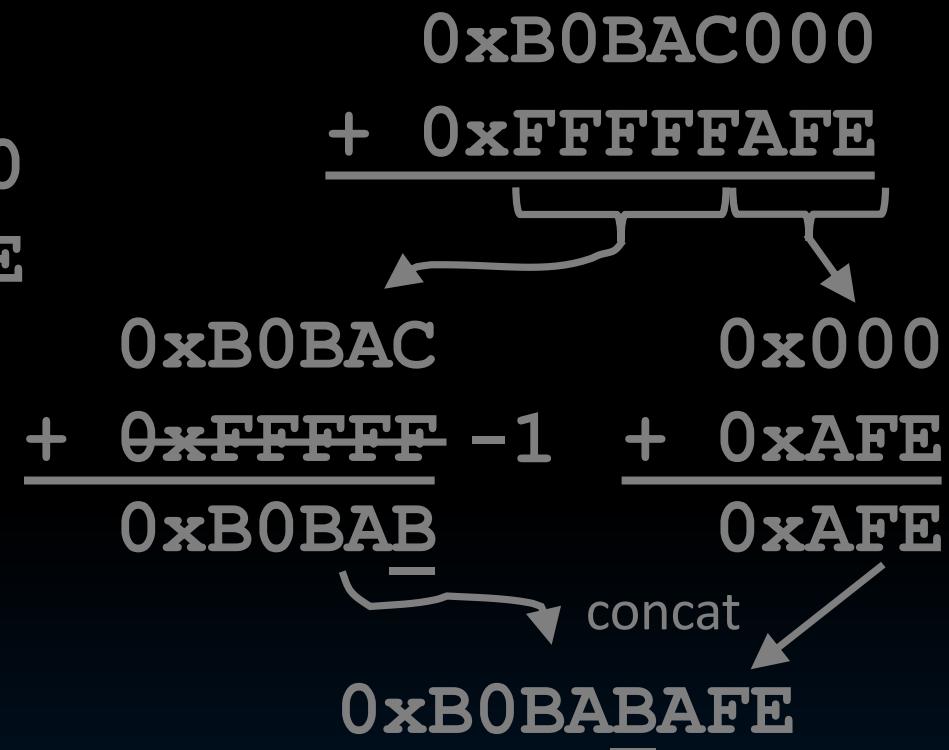


Unfortunately:

```
lui x10, 0xB0BAC      # x10 = 0xB0BAC000
addi x10, x10, 0xAF   # x10 = 0xB0BABAFE
```

Recall: **addi** sign-extends the 12-bit immediate.

If “sign bit” set, subtracts 1 from the upper 20-bits!



Solution: If 12-bit immediate is negative,  
add 1 to the upper 20-bit load.

```
lui x10, 0xB0BAD      # x10 = 0xB0BAD000
addi x10, x10, 0xAF      # x10 = 0xB0BACAFE
```

Use pseudoinstructions!  
**li** automatically handles this edge case.

# auipc loads the PC into the Register File

**auipc rd, immed**

- The **auipc** instruction
  - Adds an Upper Immediate to the PC.
- Example:

**auipc x5, 0xABCD** #  $x5 = PC + 0xABCD000$

- In Practice:

Label: **auipc x5, 0** # puts address of Label in x5

Loads the PC into a register accessible by other instructions.

**auipc** is most often used together with **jalr** to do PC-relative addressing with super large offsets.

# jalr: I-Format

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- **jalr: I-Format**

# The jalr instruction and jr pseudoinstruction

- **jalr does a Jump And Link Register:**

**jalr rd,rs1,imm**

Jump to **rs1 + imm**.

Write address of the *following* instruction to **rd**.

$$\left. \begin{array}{l} \text{PC} = \text{rs1} + \text{imm} \\ \text{rd} = \text{PC} + 4 \end{array} \right\}$$

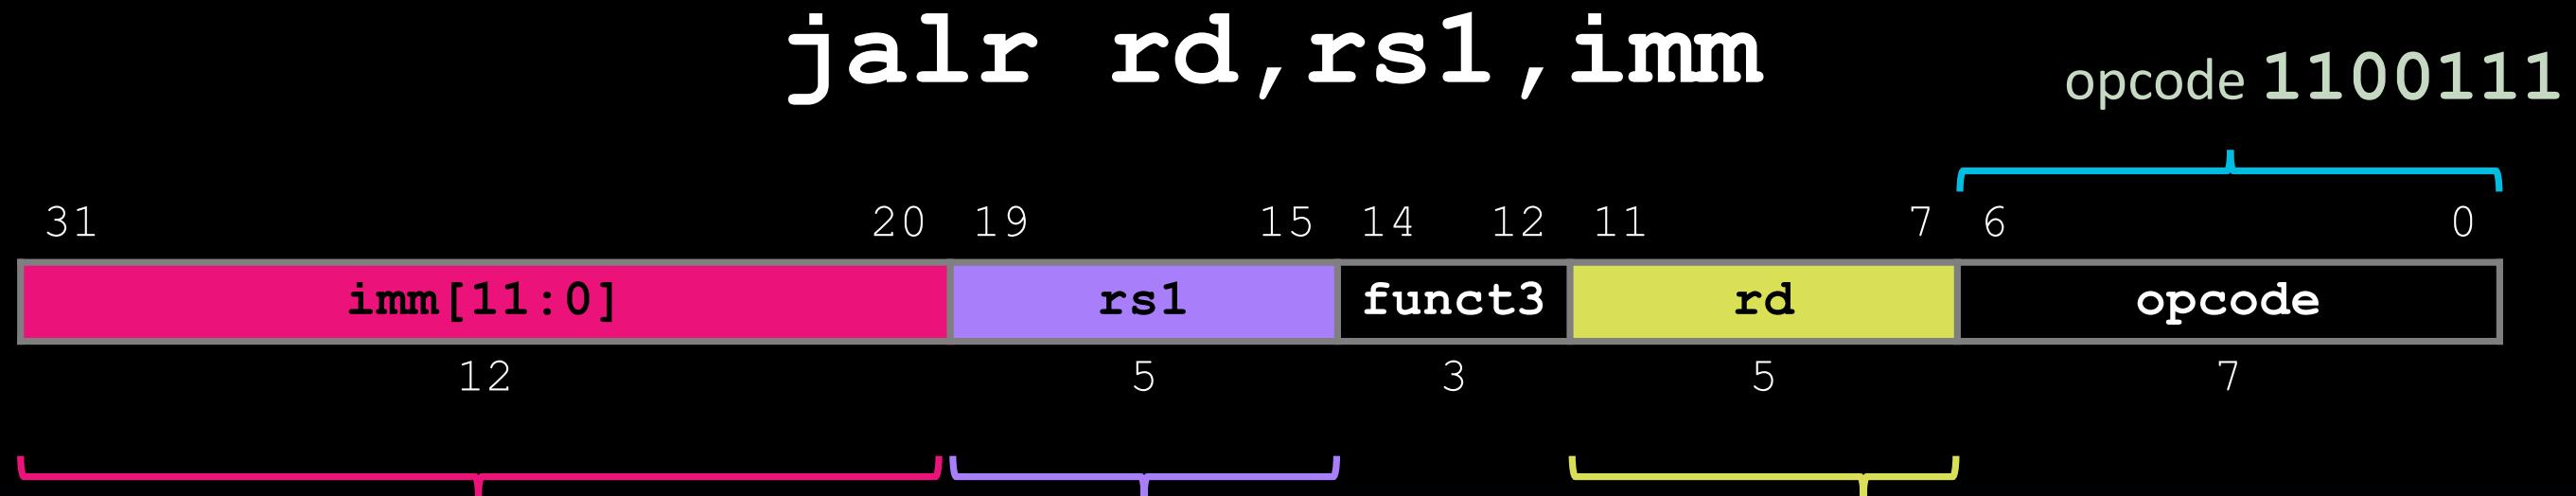
- Use cases (so far):

1. Return to caller

**jr ra**  **jalr x0, ra, 0**

# I-Format Instruction Layout: jalr

- jalr uses I-Format:



Immediate

imm and rs1 are added together to update PC.

$$PC = rs1 + imm$$

⚠ Note I-Type!

Unlike B-type, J-type, jalr will not multiply imm by 2.

Immediate offset therefore must be written in *units of bytes*.

“Destination” Register  
gets “return address.”

$$rd = PC + 4$$

# The `jalr` instruction and `jr` pseudoinstruction

- `jalr` does a Jump And Link Register:

`jalr rd, rs1, imm`

Jump to `rs1 + imm`.

Write address of the *following* instruction to `rd`.

$$\left. \begin{array}{l} \text{PC} = \text{rs1} + \text{imm} \\ \text{rd} = \text{PC} + 4 \end{array} \right\}$$

- Use cases:

1. Return to caller

`jr ra`  `jalr x0, ra, 0`

Unlike `jal` (relative to PC), `jalr` addresses are relative to `rs1`, which is modifiable by arithmetic instructions. We can do bigger jumps!

2. Call a function (and save return address)  
at any 32-bit *absolute address*.
3. Jump *PC-relative* with a *32-bit offset*.

$$\left. \begin{array}{l} \text{lui x1 <hi20bits>} \\ \text{jalr ra, x1, <lo12bits>} \\ \\ \text{auipc x1 <hi20bits>} \\ \text{jalr x0, x1, <lo12bits>} \end{array} \right\}$$

# “And in Conclusion...”

- We've covered (almost) the entire RV32 ISA!
- Practice assembling and disassembling!

imm[31:12]		rd	0110111	LUI
imm[31:12]		rd	0010111	AUIPC
imm[20 10:1 11 19:12]		rd	1101111	JAL
imm[11:0]	rs1	000	rd	JALR
imm[12 10:5]	rs2	rs1	000	BEQ
imm[12 10:5]	rs2	rs1	001	BNE
imm[12 10:5]	rs2	rs1	100	BLT
imm[12 10:5]	rs2	rs1	101	BGE
imm[12 10:5]	rs2	rs1	110	BLTU
imm[12 10:5]	rs2	rs1	111	BGEU
imm[11:0]		rs1	000	LB
imm[11:0]		rs1	001	LH
imm[11:0]		rs1	010	LW
imm[11:0]		rs1	100	LBU
imm[11:0]		rs1	101	LHU
imm[11:5]	rs2	rs1	000	SB
imm[11:5]	rs2	rs1	001	SH
imm[11:5]	rs2	rs1	010	SW
imm[11:0]		rs1	000	ADDI
imm[11:0]		rs1	010	SLTI
imm[11:0]		rs1	011	SLTIU
imm[11:0]		rs1	100	XORI
imm[11:0]		rs1	110	ORI
imm[11:0]		rs1	111	ANDI

0000000	shamt	rs1	001	rd	0010011	SLLI
0000000	shamt	rs1	101	rd	0010011	SRLI
0100000	shamt	rs1	101	rd	0010011	SRAI
0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND
0000	pred	succ	00000	000	00000	FENCE
0000	0000	0000	00000	001	00000	FENCE.I
0000000000000000			00000	000	00000	ECALL
0000000000000001			00000	000	00000	EBREAK
csr	rs1	001	rd	1110011	CSRRW	
csr	rs1	010	rd	1110011	CSRRS	
csr	rs1	011	rd	1110011	CSRRC	
csr	zimm	101	rd	1110011	CSRRWI	
csr	zimm	110	rd	1110011	CSRRSI	
csr	zimm	111	rd	1110011	CSRRCI	

Not in CS61C

<https://riscv.org/technical/specifications/>